

# Comparing different models for inferring sequency type from metadata

## Table of contents

<b>Introduction</b>	<b>1</b>
<b>Results</b>	<b>3</b>
Predictive performance . . . . .	3
Confusion matrices . . . . .	3
AUC . . . . .	20
Other metrics (precision, recall, F1-score) . . . . .	22
Feature importance . . . . .	29
Within fold consistency . . . . .	30
<b>Discussion</b>	<b>36</b>

## Introduction

Here we compare different models (and preprocessing strategies for working with mostly text-based metadata). Particularly, we compare the two (three) following strategies:

1. Textual (counting words in different columns) and numeric (characterising character-separated numeric columns as their size, minimum, maximum, sum and mean) preprocessing. We then use these values with a random forest classifier or with an XGBoost classifier
2. Plugging the raw data into a CatBoost model.

In terms of development, strategy 1. is characterized by a faster training, but a greater involvement of textual processing. Strategy 2., on the other hand, is characterised by lengthier training but a much easier development as it requires no preprocessing or hyperparameter optimization.

We also test how removing certain features - such as series description (SR), percent phase field of view (% phase FOV) and SAR - affects performance. The reason why we do this is tied with the facts that oftentimes the SR is highly specific of the centre conducting each scan, and that the % phase FOV and SAR are frequently missing from ADC sequences.

In any case we compare results hailing from cross-validation and from a hold-out test set (corresponding to 20% of the full dataset).

```
library(tidyverse)
library(knitr)

dir.create("figures",showWarnings=F)

exclusion_match <- c(
  `standard` = "All features",
  `series_description` = "No SR",
  `percent_phase_field_of_view:sar` = "No % phase FOV or SAR",
  `percent_phase_field_of_view:sar:series_description` = "No % phase FOV, SAR or SR"
)

all_metrics <- read_csv("../data_output/metrics.csv") %>%
  mutate(exclusion = factor(exclusion,
                           levels=names(exclusion_match),
                           labels=exclusion_match)) %>%

  mutate(model = factor(
    model,
    levels = c("rf","extra_trees","xgb","catboost"),
    labels = c("Random forest","Extra trees","XGBoost","CatBoost"))) %>%
  group_by(model,exclusion,metric,set) %>%
  mutate(best_fold = ifelse(
    split == "test",
    fold == fold[which.max(value[split == "cv"])],
    NA)) %>%
  subset(metric != "support") %>%
  mutate(metric = factor(
    metric,
    levels = c("auc","cm",
               "precision","recall","f1-score"),
    labels = c("auc","cm",
               "Precision","Recall","F1-score"))) %>%
  filter(set != "weighted avg") %>%
  mutate(
    set = factor(set,
```

```

      levels = c("adc","dce","dwi",
                 "t2","others",
                 "macro avg"),
      labels = c("ADC","DCE","DWI",
                 "T2W","Others",
                 "Average"))
    )
all_metrics_cv <- all_metrics %>%
  subset(split == "cv")
all_metrics_test <- all_metrics %>%
  subset(split == "test")
all_metrics_test_consensus <- all_metrics %>%
  subset(split == "test_consensus")
all_metrics_test_ensemble <- all_metrics %>%
  subset(split == "test_ensemble")

```

## Results

### Predictive performance

#### Confusion matrices

#### CV

#### As continuous variables

```

all_metrics_cv %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
  mutate(p = value / sum(value)) %>%
  group_by(model,exclusion,true,pred) %>%
  summarise(p = mean(p),
            value = mean(value)) %>%
  ggplot(aes(x = true,y = pred,
             label = sprintf("%.1f%%\n(%.0f)",p*100,value),
             fill = p)) +
  geom_tile() +
  geom_text(size = 2) +
  facet_grid(exclusion ~ model) +

```

```

theme_minimal(base_size = 8) +
scale_y_discrete(limits=rev) +
scale_fill_gradient2(low="lightskyblue1",
                     mid="white",
                     high="goldenrod1",
                     midpoint=0.5,
                     limits = c(0,1),
                     name = "") +

xlab("True") +
ylab("Predicted") +
theme(legend.position = "bottom",
      legend.key.height = unit(0.4,"cm"))
ggsave(filename="figures/cm_cv.png",
        height=7,width=7)

```




Figure 1: CV confusion matrices.

### As binned variables (0.02 intervals)

```
all_metrics_cv %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
```

```

mutate(p = value / sum(value)) %>%
group_by(model,exclusion,true,pred) %>%
summarise(p = mean(p),
           value = mean(value)) %>%
mutate(p_fill = cut(
  p,seq(0,1,by=0.02),
  include.lowest=T)) %>%
ggplot(aes(x = true,y = pred,
           label = sprintf("%.1f%%\n(%.0f)",p*100,value),
           fill = p_fill)) +
geom_tile(alpha = 0.6) +
geom_text(size = 2) +
facet_grid(exclusion ~ model) +
theme_minimal(base_size = 8) +
scale_y_discrete(limits=rev) +
scale_fill_brewer(palette = "PiYG",name = "") +
xlab("True") +
ylab("Predicted") +
theme(legend.position = "bottom",
      legend.key.height = unit(0.4,"cm"),
      panel.grid = element_blank())
ggsave(filename="figures/cm_cat_cv.png",
        height=7,width=7)

```



Figure 2: CV confusion matrices (binned).

**Hold-out test-set**

**As continuous variables**

```

all_metrics_test %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
  mutate(p = value / sum(value)) %>%
  group_by(model,exclusion,true,pred) %>%
  summarise(p = mean(p),
            value = median(value)) %>%
  ggplot(aes(x = true,y = pred,
            label = sprintf("%.1f%%\n(%s)",p*100,value),
            fill = p)) +
  geom_tile() +
  geom_text(size = 2) +
  facet_grid(exclusion ~ model) +
  theme_minimal(base_size = 8) +
  scale_y_discrete(limits=rev) +
  scale_fill_gradient2(low="lightskyblue1",
                      mid="white",
                      high="goldenrod1",
                      midpoint=0.5,
                      limits = c(0,1),
                      name = "") +

  xlab("True") +
  ylab("Predicted") +
  theme(legend.position = "bottom",
        legend.key.height = unit(0.4,"cm"))
ggsave(filename="figures/cm_test.png",
        height=7,width=7)

```






Figure 3: Hold-out test-set confusion matrices.

### As binned variables (0.02 intervals)

```
all_metrics_test %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
```

```

mutate(p = value / sum(value)) %>%
group_by(model,exclusion,true,pred) %>%
summarise(p = mean(p),
           value = median(value)) %>%
mutate(p_fill = cut(p,seq(0,1,by=0.02),
                     include.lowest=T)) %>%
ggplot(aes(x = true,y = pred,
           label = sprintf("%.1f%%\n(%.0f)",p*100,value),
           fill = p_fill)) +
geom_tile(alpha = 0.6) +
geom_text(size = 2) +
facet_grid(exclusion ~ model) +
theme_minimal(base_size = 8) +
scale_y_discrete(limits=rev) +
scale_fill_brewer(palette = "PiYG",name = "") +
xlab("True") +
ylab("Predicted") +
theme(legend.position = "bottom",
      legend.key.height = unit(0.4,"cm"),
      panel.grid = element_blank())
ggsave(filename="figures/cm_cat_test.png",
        height=7,width=7)

```



Figure 4: Hold-out test-set confusion matrices (binned).

**Hold-out test-set (consensus)**

**As continuous variables**

```

all_metrics_test_consensus %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
  mutate(p = value / sum(value)) %>%
  ggplot(aes(x = true,y = pred,
             label = sprintf("%.1f%%\n(%s)",p*100,value),
             fill = p)) +
  geom_tile() +
  geom_text(size = 2) +
  facet_grid(exclusion ~ model) +
  theme_minimal(base_size = 8) +
  scale_y_discrete(limits=rev) +
  scale_fill_gradient2(low="lightskyblue1",
                      mid="white",
                      high="goldenrod1",
                      midpoint=0.5,
                      limits = c(0,1),
                      name = "") +

  xlab("True") +
  ylab("Predicted") +
  theme(legend.position = "bottom",
        legend.key.height = unit(0.4,"cm"))
ggsave(filename="figures/cm_test_consensus.png",
        height=7,width=7)

```




Figure 5: Hold-out test-set consensus confusion matrices.

### As binned variables (0.02 intervals)

```
all_metrics_test_consensus %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
```

```

mutate(p = value / sum(value)) %>%
mutate(p_fill = cut(p,seq(0,1,by=0.02),
                      include.lowest=T)) %>%
ggplot(aes(x = true,y = pred,
           label = sprintf("%.1f%%\n(%.0f)",p*100,value),
           fill = p_fill)) +
geom_tile(alpha = 0.6) +
geom_text(size = 2) +
facet_grid(exclusion ~ model) +
theme_minimal(base_size = 8) +
scale_y_discrete(limits=rev) +
scale_fill_brewer(palette = "PiYG",name = "") +
xlab("True") +
ylab("Predicted") +
theme(legend.position = "bottom",
      legend.key.height = unit(0.4,"cm"),
      panel.grid = element_blank())
ggsave(filename="figures/cm_cat_test_consensus.png",
        height=7,width=7)

```



Figure 6: Hold-out test-set consensus confusion matrices (binned).

**Hold-out test-set (ensemble)**

**As continuous variables**

```

all_metrics_test_ensemble %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
  mutate(p = value / sum(value)) %>%
  ggplot(aes(x = true,y = pred,
             label = sprintf("%.1f%%\n(%s)",p*100,value),
             fill = p)) +
  geom_tile() +
  geom_text(size = 2) +
  facet_grid(exclusion ~ model) +
  theme_minimal(base_size = 8) +
  scale_y_discrete(limits=rev) +
  scale_fill_gradient2(low="lightskyblue1",
                      mid="white",
                      high="goldenrod1",
                      midpoint=0.5,
                      limits = c(0,1),
                      name = "") +

  xlab("True") +
  ylab("Predicted") +
  theme(legend.position = "bottom",
        legend.key.height = unit(0.4,"cm"))
ggsave(filename="figures/cm_test_ensemble.png",
        height=7,width=2.5)

```






Figure 7: Hold-out test-set ensemble confusion matrices.

As binned variables (0.02 intervals)

```

all_metrics_test_ensemble %>%
  subset(metric == "cm") %>%
  group_by(model,exclusion,true,fold) %>%
  mutate(p = value / sum(value)) %>%
  mutate(p_fill = cut(p,seq(0,1,by=0.02),
                      include.lowest=T)) %>%
  ggplot(aes(x = true,y = pred,
             label = sprintf("%.1f%%\n(%.0f)",p*100,value),
             fill = p_fill)) +
  geom_tile(alpha = 0.6) +
  geom_text(size = 2) +
  facet_grid(exclusion ~ model) +
  theme_minimal(base_size = 8) +
  scale_y_discrete(limits=rev) +
  scale_fill_brewer(palette = "PiYG",name = "") +
  xlab("True") +
  ylab("Predicted") +
  theme(legend.position = "bottom",
        legend.key.height = unit(0.4,"cm"),
        panel.grid = element_blank())
ggsave(filename="figures/cm_test_ensemble.png",
        height=7,width=2.5)

```

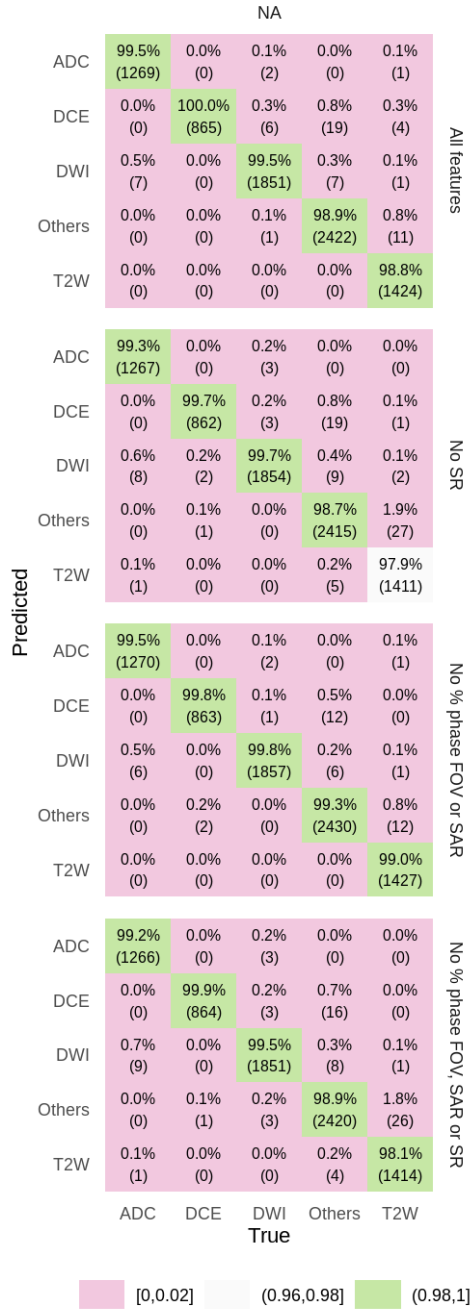


Figure 8: Hold-out test-set ensemble confusion matrices (binned).

## AUC

## CV

```
all_metrics_cv %>%
  subset(metric == "auc") %>%
  group_by(model,exclusion) %>%
  summarise(vmin = min(value),
            vmax = max(value),
            value = mean(value)) %>%
  ggplot(aes(y = value,x = model,
            ymin = vmin,ymax=vmax,
            label = sprintf("%.2f%%",value*100),
            colour = exclusion)) +
  geom_point(position = position_dodge(0.8)) +
  geom_linerange(position = position_dodge(0.8)) +
  theme_minimal(base_size = 8) +
  scale_colour_brewer(type = "qual",palette = 2,
                    name = "Feature subset") +
  xlab("Models") +
  ylab("AUC") +
  scale_x_discrete(limits = rev) +
  scale_y_continuous(
    labels = function(x) sprintf("%.2f%%",x*100)) +
  theme(legend.key.height = unit(0.2,"cm"),
        legend.key.width = unit(0.2,"cm")) +
  coord_flip() +
  theme(legend.position = "bottom",
        legend.key.height = unit(0.4,"cm")) +
  guides(colour = guide_legend(nrow = 2))
ggsave(filename="figures/auc_cv.png",
        height=2,width=4)
```

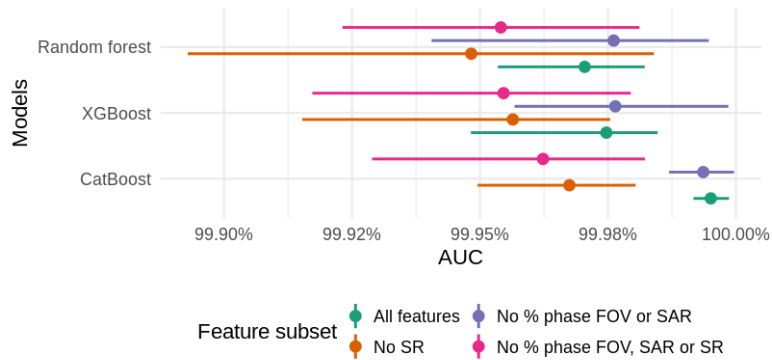


Figure 9: CV AUC.

## Hold-out test set

```
all_metrics_test %>%
  subset(metric == "auc") %>%
  group_by(model,exclusion) %>%
  summarise(vmin = min(value),
            vmax = max(value),
            value = mean(value)) %>%
  ggplot(aes(y = value,x = model,
            ymin = vmin,ymax=vmax,
            label = sprintf("%.2f%%",value*100),
            colour = exclusion)) +
  geom_point(position = position_dodge(0.8)) +
  geom_linerange(position = position_dodge(0.8)) +
  geom_point(
    data=subset(all_metrics_test,
                metric == "auc" & best_fold == T),
    aes(x = model,y = value,
        group = exclusion,
        shape = ifelse(best_fold,"Best CV model")),
    inherit.aes=F,
    position = position_dodge(0.8),
    size = 4) +
  scale_shape_manual(values = c("+"),
                    guide = F) +
  theme_minimal(base_size = 8) +
  xlab("Models") +
  ylab("AUC") +
```

```

scale_colour_brewer(type = "qual", palette = 2,
                    name = "Feature subset") +
scale_x_discrete(limits = rev) +
scale_y_continuous(
  labels = function(x) sprintf("%.2f%%", x*100)) +
theme(legend.key.height = unit(0.2, "cm"),
      legend.key.width = unit(0.2, "cm")) +
coord_flip() +
theme(legend.position = "bottom",
      legend.key.height = unit(0.4, "cm")) +
guides(colour = guide_legend(nrow = 2))
ggsave(filename="figures/auc_test.png",
        height=2,width=4)

```

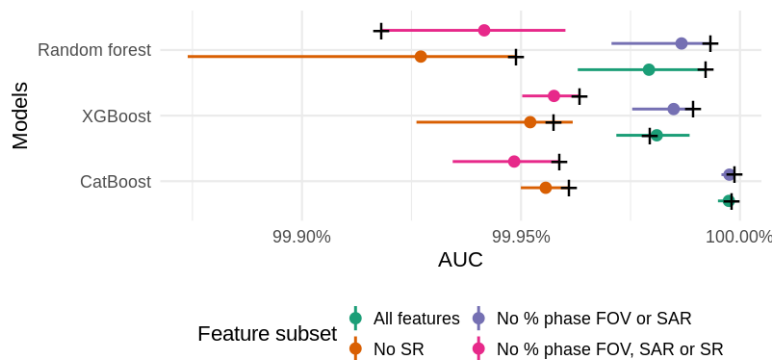


Figure 10: Hold-out test-set AUC.

## Other metrics (precision, recall, F1-score)

### CV

```

M <- c("Precision", "Recall", "F1-score")

all_metrics_cv %>%
  subset(metric %in% M) %>%
  group_by(model, exclusion, metric, set) %>%
  summarise(vmin = min(value),
            vmax = max(value),
            value = mean(value)) %>%
  ggplot(aes(y = value, x = model,

```

```

      ymin = vmin,ymax=vmax,
      label = sprintf("%.2f%%",value*100),
      colour = exclusion)) +
geom_point(position = position_dodge(0.8)) +
geom_linerange(position = position_dodge(0.8)) +
theme_minimal(base_size = 8) +
xlab("Models") +
ylab("Value") +
facet_grid(set ~ metric) +
scale_colour_brewer(type = "qual",palette = 2,
                    name = "Feature subset") +
scale_y_continuous(
  labels = function(x) sprintf("%.2f%%",x*100)) +
theme(legend.key.height = unit(0.2,"cm"),
      legend.key.width = unit(0.2,"cm"),
      panel.background = element_rect(
        fill = NA,colour = "black")) +
coord_flip() +
theme(legend.position = "bottom",
      legend.key.height = unit(0.4,"cm")) +
guides(colour = guide_legend(nrow = 2))
ggsave(filename="figures/metrics_cv.png",
        height=7,width=7)

```




Figure 11: CV precision, recall and F1-score.

### Hold-out test-set

```
all_metrics_test %>%
  subset(metric %in% M) %>%
  group_by(model,exclusion,metric,set) %>%
```



```

summarise(vmin = min(value),
          vmax = max(value),
          value = mean(value)) %>%
ggplot(aes(y = value,x = model,
          ymin = vmin,ymax=vmax,
          label = sprintf("%.2f%%",value*100),
          colour = exclusion)) +
geom_point(position = position_dodge(0.8)) +
geom_linerange(position = position_dodge(0.8)) +
geom_point(
  data=subset(all_metrics_test,
              metric %in% M & best_fold == T),
  aes(x = model,y = value,
      group = exclusion,
      shape = ifelse(best_fold,"Best CV model")),
  inherit.aes=F,
  position = position_dodge(0.8),
  size = 3) +
scale_shape_manual(values = c("+"),
                   guide = F) +
theme_minimal(base_size = 8) +
xlab("Models") +
ylab("Value") +
facet_grid(set ~ metric) +
scale_colour_brewer(type = "qual",palette = 2,
                   name = "Feature subset") +
scale_y_continuous(
  labels = function(x) sprintf("%.2f%%",x*100)) +
theme(legend.key.height = unit(0.2,"cm"),
      legend.key.width = unit(0.2,"cm"),
      panel.background = element_rect(
        fill = NA,colour = "black")) +
coord_flip() +
theme(legend.position = "bottom",
      legend.key.height = unit(0.4,"cm")) +
guides(colour = guide_legend(nrow = 2))
ggsave(filename="figures/metrics_test.png",
        height=7,width=7)

```




Figure 12: Hold-out test-set precision, recall and F1-score.

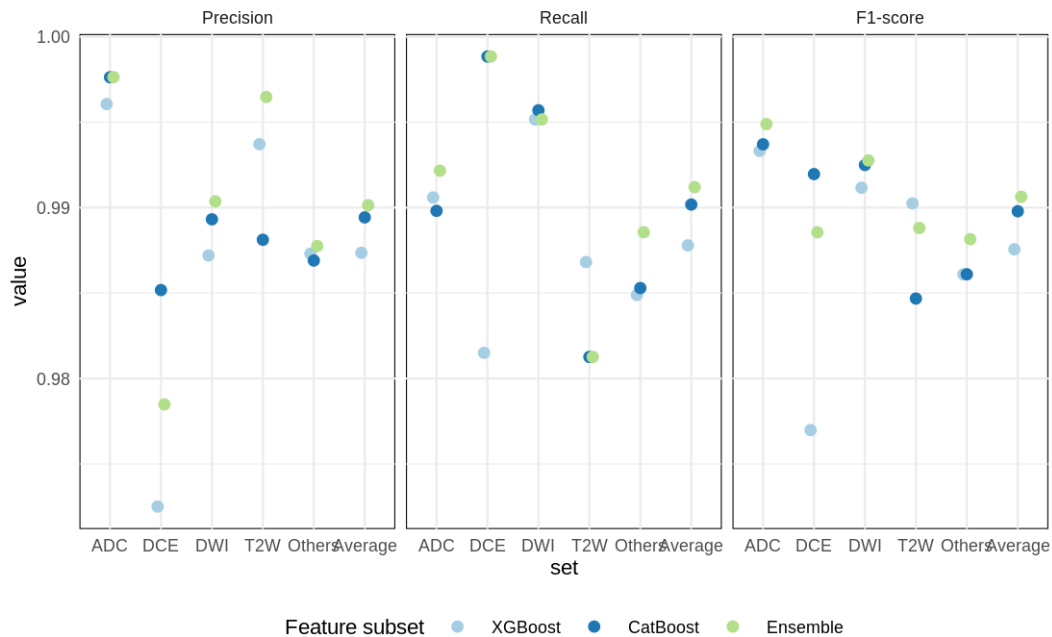
```
tmp_df <- all_metrics_test_ensemble %>%
  subset(metric %in% M) %>%
  subset(exclusion == "No % phase FOV, SAR or SR") %>%
  mutate(model = "Ensemble")
```

```

tmp_df <- all_metrics_test_consensus %>%
  subset(metric %in% M) %>%
  subset(model %in% c("XGBoost","CatBoost") &
    exclusion == "No % phase FOV, SAR or SR") %>%
  rbind(tmp_df) %>%
  mutate(model = factor(
    model,
    levels = c("XGBoost","CatBoost","Ensemble")))

tmp_df %>%
  ggplot(aes(x = set,y = value,colour = model)) +
  geom_point(position=position_dodge(width=0.2)) +
  facet_grid(~ metric) +
  theme_minimal(base_size = 8) +
  scale_colour_brewer(type = "qual",palette = 3,
    name = "Feature subset") +
  theme(legend.position = "bottom",
    legend.key.height = unit(0.4,"cm"),
    panel.background = element_rect(
      fill=NA,colour="black"))

```



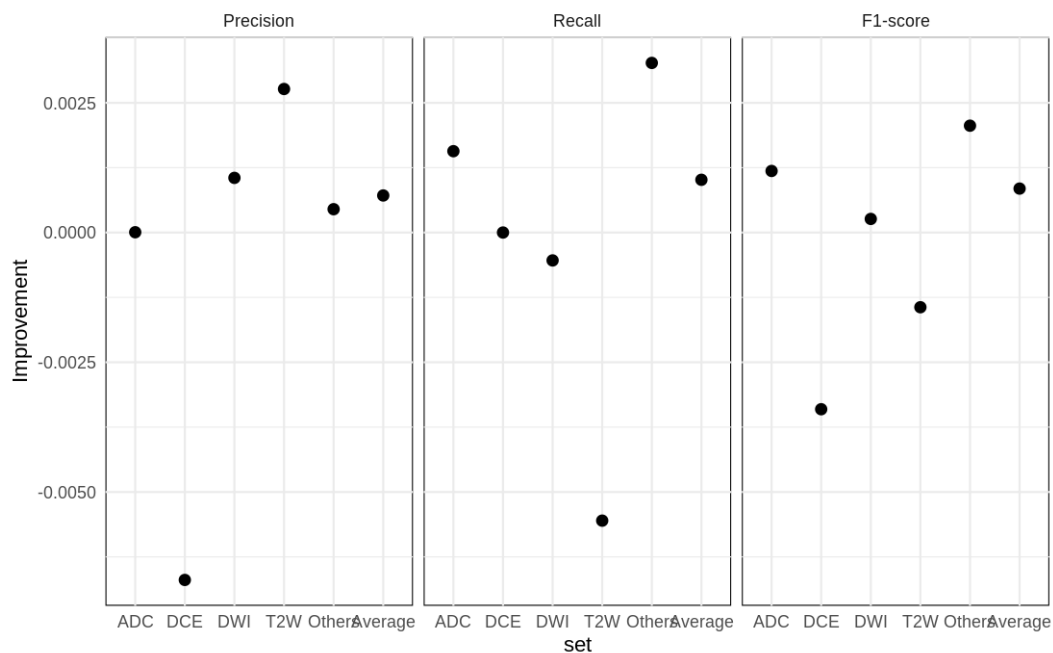
```
tmp_df %>%
  group_by(metric,set) %>%
  summarise(Improvement = value[model == "Ensemble"] - max(value[model != "Ensemble"]))
  group_by(metric) %>%
  subset(set != "Average") %>%
  summarise(`Average Improvement` = mean(Improvement))
```

`summarise()` has grouped output by 'metric'. You can override using the  
`.groups` argument.

```
# A tibble: 3 x 2
  metric      `Average Improvement`
  <fct>          <dbl>
1 Precision      -0.000484
2 Recall         -0.000251
3 F1-score       -0.000267
```

```
tmp_df %>%
  group_by(metric,set) %>%
  summarise(Improvement = value[model == "Ensemble"] - max(value[model != "Ensemble"]))
  ggplot(aes(x = set,y = Improvement)) +
  geom_point(position=position_dodge(width=0.2)) +
  facet_grid(~ metric) +
  theme_minimal(base_size = 8) +
  scale_colour_brewer(type = "qual",palette = 3,
    name = "Feature subset") +
  theme(legend.position = "bottom",
    legend.key.height = unit(0.4,"cm"),
    panel.background = element_rect(
      fill=NA,colour="black"))
```

`summarise()` has grouped output by 'metric'. You can override using the  
`.groups` argument.



## Feature importance

```
feature_importance <- read_csv("../data_output/feature_importances.csv") %>%
  mutate(sub_feature = str_match(feature, ".*"),
         feature = ifelse(
           grepl(":", feature),
           gsub(":", "", str_match(feature, ".*:")),
           feature)) %>%
  group_by(model, feature, class, fold, exclusion) %>%
  summarise(value = value[which.max(abs(value))]) %>%
  mutate(model = factor(
    model,
    levels = c("rf", "extra_trees", "xgb", "catboost"),
    labels = c("Random forest", "Extra trees", "XGBoost", "CatBoost"))) %>%
  mutate(exclusion = factor(exclusion,
                           levels=names(exclusion_match),
                           labels=exclusion_match))
```

Rows: 34870 Columns: 6

-- Column specification -----

Delimiter: ","

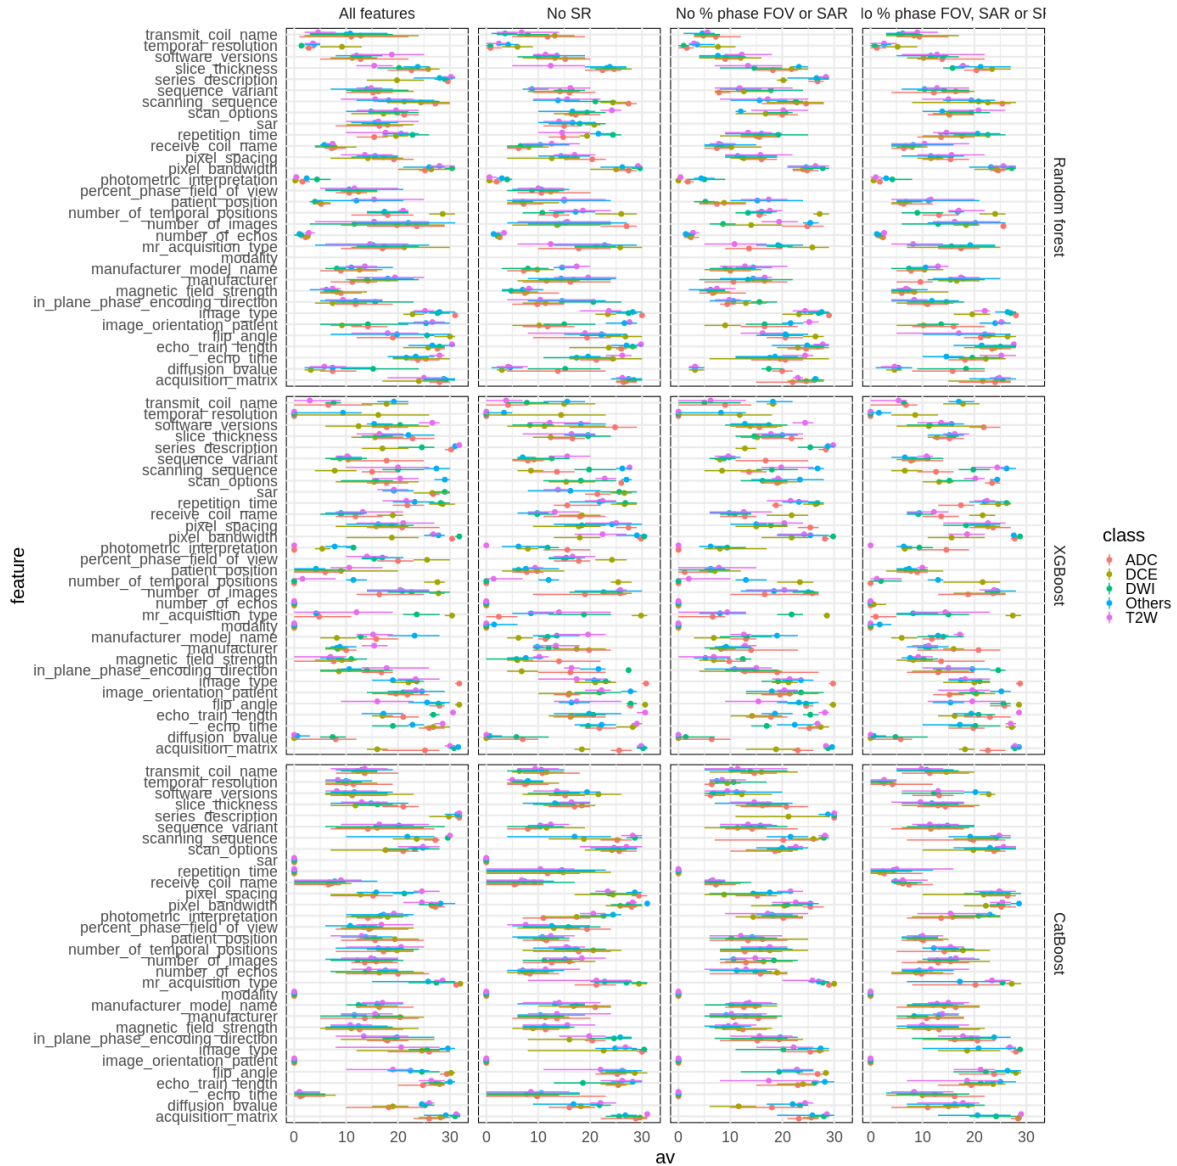
```
chr (4): model, exclusion, class, feature
dbl (2): value, fold
```

i Use ``spec()`` to retrieve the full column specification for this data.  
i Specify the column types or set ``show_col_types = FALSE`` to quiet this message.  
``summarise()`` has grouped output by 'model', 'feature', 'class', 'fold'. You can override using ``ungroup()``.

## Within fold consistency

```
feature_importance %>%
  group_by(model, fold, exclusion, class) %>%
  mutate(R = sign(value) * rank(abs(value))) %>%
  group_by(model, feature, exclusion, class) %>%
  summarise(
    av = mean(abs(R)), m = min(abs(R)), M = max(abs(R)) %>%
  ggplot(aes(feature, y = av, ymin = m, ymax = M, colour = class)) +
  geom_point(size = 0.5, position = position_dodge(0.5)) +
  geom_linerange(size = 0.25, position = position_dodge(0.5)) +
  facet_grid(model ~ exclusion) +
  coord_flip() +
  theme_minimal(base_size = 8) +
  theme(legend.key.height = unit(0.2, "cm"),
        legend.key.width = unit(0.2, "cm"),
        panel.background = element_rect(
          fill = NA, colour = "black"))
```

``summarise()`` has grouped output by 'model', 'feature', 'exclusion'. You can override using the ``.groups`` argument.



```
ggsave(filename="figures/feature_importance.png",
        height=7,width=7)
```

```
library(cowplot)
```

```
best_folds <- all_metrics %>%
  filter(metric == "auc" & split == "cv") %>%
```

```

subset(best_fold) %>%
ungroup %>%
select(model,exclusion,fold) %>%
distinct

plot_df <- merge(
  feature_importance,best_folds,
  by = c("model","exclusion","fold"),all=F) %>%
group_by(model,exclusion,class) %>%
mutate(R = sign(value) * rank(abs(value))) %>%
group_by(model,feature,exclusion,class) %>%
summarise(value = mean(value),
           average_rank = mean(R)) %>%
group_by(model,feature,exclusion) %>%
mutate(order = sum(abs(value)))

```

`summarise()` has grouped output by 'model', 'feature', 'exclusion'. You can override using the `.groups` argument.

```

all_plots <- list()

for (m in unique(plot_df$model)) {
  for (exc in unique(plot_df$exclusion)) {
    str_ <- sprintf("%s %s",m,exc)
    all_plots[[str_]] <- plot_df %>%
      subset(model == m & exclusion == exc) %>%
      ggplot(aes(reorder(feature,order),
                  y = value,
                  colour = class)) +
      geom_point(size = 0.5) +
      coord_flip() +
      theme_minimal(base_size = 8) +
      theme(legend.key.height = unit(0.2,"cm"),
            legend.key.width = unit(0.2,"cm"),
            panel.background = element_rect(
              fill = NA,colour = "black"),
            legend.position = "bottom") +
      ggtitle(m,subtitle=exc) +
      xlab("") +
      ylab("SHAP value") +
      scale_colour_brewer(

```



```
        type = "qual",palette = 3,  
        name = NULL)  
    }  
}  
  
plot_grid(plotlist=all_plots[1:4])
```

```
ggsave(  
  filename="figures/feature_importance_best_fold_1.png",  
  height=6,width=8)  
  
plot_grid(plotlist=all_plots[5:8])
```

```
ggsave(  
  filename="figures/feature_importance_best_fold_2.png",  
  height=6,width=8)  
  
plot_grid(plotlist=all_plots[9:12])
```

```
ggsave(  
  filename="figures/feature_importance_best_fold_3.png",  
  height=6,width=8)
```

## Discussion

The consensus is relatively easy to reach - CatBoost is, by far, the best performing model out of the two. The advantages are considerable - the optimization is simpler and no preprocessing is required. While it takes longer to train, one can be sure that the results will be hard to dispute.